

AWS · SAAS PRICING AND PACKAGING

SaaS pricing and packaging on AWS

From business model design to technical execution — and the agentic AI frontier





SESSION OVERVIEW

What we'll cover

01 The pricing opportunity

03 The CFO view: unit economics

05 The agentic AI frontier

02 Packaging for growth

04 Technical foundations on AWS

CHAPTER 01

The pricing opportunity



BY THE NUMBERS

2–4×

SaaS companies that actively manage pricing strategy grow 2–4× faster than those that set prices once and leave them. Pricing is the highest-leverage growth decision you make.

Source: AWS — Monetization Strategies for SaaS on AWS



PRICING STRATEGY

Three decisions every SaaS leader must get right

- | | | |
|----|-------------------------------|--|
| 01 | What to charge for | Choose a value metric that scales with customer success — seats, API calls, data processed, outcomes delivered. Misalignment here creates churn. |
| 02 | How to structure tiers | Good/Better/Best packaging guides buyers to the right plan and protects margin. Each tier should have a clear buyer persona and upgrade trigger. |
| 03 | How to go to market | Self-serve, sales-assisted, or marketplace-led. Your GTM motion must match your buyers' procurement preferences — especially in enterprise. |



PRICING MODELS

Five dominant SaaS pricing approaches



Flat rate

Single price for full access. Simple to buy, simple to sell. Hard to capture value at scale.

E.G. \$499/MONTH ALL-IN



Per seat

Price per user. Predictable, easy to budget. Revenue grows with adoption — but can discourage broad rollout.

E.G. \$30/USER/MONTH



Usage-based

Pay for what you consume. Lowers entry barrier, aligns cost to value. The dominant model for AI and data workloads.

E.G. \$0.002/API CALL



Tiered

Price changes at usage thresholds. Rewards high-volume customers, smooths expansion revenue curve.

E.G. FIRST 10K CALLS FREE, THEN \$0.001



Feature-based

Gate premium capabilities behind higher plans. Drives upsell through feature discovery and business need.

E.G. SSO, AUDIT LOGS ON ENTERPRISE

CHAPTER 02

Packaging for growth



PACKAGING ARCHITECTURE

Good, Better, Best — the proven packaging structure

STARTER

Good

Self-serve entry point

- ✓ Core product value
- ✓ Usage limits to control cost
- ✓ Email support
- ✓ Advanced features
- ✓ Enterprise security

MOST POPULAR

Better

Where most revenue lands

- ✓ Everything in Starter
- ✓ Higher usage limits
- ✓ Priority support + SLA
- ✓ Advanced analytics
- ✓ Custom contracts

ENTERPRISE

Best

Negotiated, high-ACV

- ✓ Everything in Better
- ✓ Unlimited or committed usage
- ✓ SSO, audit logs, SCIM
- ✓ Dedicated CSM
- ✓ Custom contracts and SLAs



Your architecture is a pricing decision

DEDICATED RESOURCES

Silo

Each tenant gets their own isolated stack

■ Tenant A — dedicated infra

■ Tenant B — dedicated infra

■ Tenant C — dedicated infra

Shared: onboarding, identity, control plane

- Easy per-tenant cost attribution — all resources scoped to one account
- Compliance-ready — hard isolation for regulated workloads
- Higher per-tenant cost — trade-off for isolation premium pricing

SHARED RESOURCES

Pool

All tenants share infrastructure

Tenant A

Tenant B

Tenant C

Shared compute, storage, messaging

Shared: onboarding, identity, control plane

- Economies of scale — lowest infrastructure cost per tenant
- Requires fine-grained metering for accurate cost attribution
- Noisy-neighbour risk — needs throttling and quota enforcement

HYBRID — THE REALITY

Bridge

Mix silo and pool by service or tier

Enterprise tenants — siloed data layer

Standard tenants — pooled compute

Shared: onboarding, identity, control plane

- Match your architecture cost to your pricing tier — charge more for silo isolation
- Regulatory profile and noisy-neighbour risk guide service-level decisions
- Most common real-world pattern — requires careful tenant context propagation

CHAPTER 03

The CFO view: unit economics



FINANCIAL HEALTH

Pricing decisions live in your income statement

NRR TARGET

>120%

Net revenue retention above 120% means your existing book grows without a single new logo. Pricing expansion drives this.

LTV:CAC RATIO

3:1

Minimum healthy ratio. Usage-based models improve LTV as customers grow into the product — CAC stays fixed.

PAYBACK PERIOD

<18mo

Time to recover CAC from gross margin. Tier packaging with a clear upgrade path compresses this meaningfully.

GROSS MARGIN


70–80%

Target for scalable SaaS. AI workloads compress this to 50–60% — fine-grained cost attribution is how you defend it.



Five levels from ad-hoc spending to real-time pricing feedback

Most SaaS companies operate at Level 2–3 · unit metrics require Level 4+ instrumentation





Unit metric = cost per unit of business value

NUMERATOR

AWS cost

From Cost and Usage Report (CUR) — tag resources by tenant, product, and environment to isolate production spend.

DENOMINATOR

Business driver

A demand metric correlated to costs: API calls, active tenants, documents processed. Sourced from CloudWatch, your database, or a data lake.

\$2.10

per 1,000 API responses

API-first SaaS — maps directly to a usage-based price of ~\$3–4/1K at 60% margin

\$4.80

per active tenant / month

Pooled model baseline — identifies margin outlier tenants consuming 10× average

\$0.03

per document analysed

Document AI — direct input for outcome-based pricing at \$0.10–0.15/doc

\$0.12

per agent task completed

Agentic AI — sets floor for workflow pricing; hybrid fee + per-task keeps margin intact



Two strategies to measure cost per tenant

LESS INVASIVE · FAST TO IMPLEMENT

Coarse-grained

- Approximate tenant cost from activity share — if Tenant A is 10% of active users, allocate 10% of charges
- Good for siloed models where resources map directly to one tenant
- Suitable for early-stage products and pooled architectures with uniform tenants
- Lacks precision — won't reveal noisy neighbours or margin outliers

COST & USAGE REPORTS

RESOURCE TAGS

CLOUDWATCH

DETAILED INSTRUMENTATION · FULL ATTRIBUTION

Fine-grained

- Instrument at application and infrastructure layer — capture API counts, DB operations, compute time per tenant
- Publish metrics as events; aggregate into per-tenant cost records in near real-time
- Enables precision pricing, margin defence, and proactive throttling of outlier tenants
- Required for usage-based and outcome-based billing at scale

CLOUDWATCH

LAMBDA

KINESIS FIREHOSE

DYNAMODB

S3


CHAPTER 04

Technical foundations on AWS



REFERENCE ARCHITECTURE

Usage metering — from API call to billing record





GO-TO-MARKET CHANNEL

AWS Marketplace removes procurement friction

330K+

active customers with pre-committed AWS spend (DRAW / EDP)

Weeks

reduction in enterprise sales cycle — procurement already approved

Why it matters for your GTM

- Buyers consume your SaaS against existing AWS committed spend — no new PO process
- Supports usage-based, annual contract, and free-trial pricing natively
- Private Offers let you negotiate custom terms without leaving the marketplace billing rail
- Co-sell with AWS field teams accelerates enterprise deals at scale



GETTING STARTED

A 90-day roadmap to execution

DAYS 1–30

Instrument and measure

Define your value metric — what usage correlates to customer success

Instrument with Kinesis + Lambda metering pipeline

Baseline current COGS per tenant using fine-grained attribution

Audit current pricing vs. competitor landscape

DAYS 31–60

Package and test

Design Good/Better/Best tier structure with clear upgrade triggers

A/B test pricing page with real prospects

Choose silo/pool/bridge model to match tier economics

Set NRR, LTV:CAC, and payback period targets

DAYS 61–90

Activate and scale

List on AWS Marketplace with usage-based billing

Enable Private Offers for enterprise deals

Co-sell activation with your AWS partner development manager

Launch expansion playbook targeting existing accounts

CHAPTER 05

The agentic AI frontier

How autonomous AI is forcing a full repricing of what SaaS delivers



Agents don't have seats. Per-seat pricing breaks.

TRADITIONAL SAAS

Human uses software

Users have seats. Usage is bounded by working hours. COGS is near-zero per seat. Per-seat pricing captures expansion through headcount.

GROSS MARGIN: 80–90%

AGENTIC AI

Agent acts autonomously

No seats — agents run 24/7, execute multi-step processes independently. Every invocation has real inference COGS: model tokens, tool calls, memory reads, execution time.

GROSS MARGIN: 50–60% WITHOUT ACTIVE COST MANAGEMENT

THE 2026 RENEWAL CLIFF

2025 AI pilots priced for adoption. 2026 production contracts must reflect delivered value. Soft ROI positioning doesn't survive renewal. Source: Bessemer Venture Partners

THE IMPLICATION

Pricing must shift from access to outcomes. The COMPASS framework (AWS + Zuora + Simon-Kucher) provides a structured approach: map scope of work against attributability to find your pricing lever.



Three charge metrics for AI products

01

Consumption

PER TOKEN · PER API CALL

Mirrors your infrastructure economics. Predictable margin, easy to model. Best for technical buyers who understand token counting.

Challenge: buyers think in outcomes, not tokens. Confusing for non-technical enterprise buyers.

MARGIN PREDICTABILITY

VALUE ALIGNMENT

02

Workflow

PER TASK · PER COMPLETED JOB

Charges for discrete, recognisable work units — book a meeting, analyse a spreadsheet, process a document. Dramatically clarifies the value proposition.

Best of both: buyers understand what they're buying, you can model costs per workflow type.

MARGIN PREDICTABILITY

VALUE ALIGNMENT

03

Outcome

PER RESOLUTION · PER RESULT

Maximum value alignment. Charges only when the AI delivers a measurable result. Highest willingness-to-pay — and highest cost variability.


Real examples: Intercom \$0.99/ticket resolved · EvenUp per legal package · Leena AI per ticket closed

MARGIN PREDICTABILITY

VALUE ALIGNMENT




Multi-tenant agent isolation mirrors your pricing tiers





Map your agent to the right pricing lever



Five principles for AI pricing success

- 01 **Metrics alignment** — choose the metric that reflects actual value delivery, not AI activity
- 02 **Unit economics first** — if margins don't work at 10 customers, they won't at 1,000
- 03 **Test, don't model** — raise price incrementally until customers say "we need to think about that"
- 04 **Hybrid bridges the gap** — platform fee (2× delivery cost) + outcome credits captures expansion without sticker shock
- 05 **Tier your architecture** — silo/pool/bridge on AgentCore should map to your pricing tiers



FURTHER READING

Sources that informed this deck

PRICING STRATEGY

The SaaS Monetization Playbook

Bessemer Venture Partners — consumption models, AI pricing, the 2026 renewal cliff

bvp.com/atlas/the-saas-monetization-playbook →

MULTI-TENANCY

Silo, pool, and bridge models

AWS Well-Architected SaaS Lens — isolation models and their pricing implications

docs.aws.amazon.com/wellarchitected/saas-lens →

UNIT METRICS

Choose, create, and track unit metrics

AWS — selecting business drivers, building a metering pipeline, visualising in QuickSight

aws.amazon.com/blogs/mt/unit-metrics-for-your-applications →

AI MONETISATION

Pricing strategies for ISVs building AI

AWS — COMPASS framework, consumption vs workflow vs outcome charge metrics

aws.amazon.com/blogs/apn/pricing-strategies-for-isv-partners →

COST ATTRIBUTION

Calculating cost per tenant

AWS — coarse-grained vs fine-grained attribution strategies in pooled architectures

aws.amazon.com/blogs/apn/calculating-tenant-costs-in-saas →

AGENTIC AI INFRASTRUCTURE

Amazon Bedrock AgentCore

AWS — multi-tenant agent isolation, silo/pool/bridge patterns for production deployments

aws.amazon.com/bedrock/agentcore →



AWS · SAAS PRICING AND PACKAGING

Pricing is a product decision. Make it with the same rigour as your architecture.

[SaaS on AWS](#) →

[AWS Marketplace](#) →

[Amazon Bedrock AgentCore](#) →